

Midterm Review

DS 5110: Big Data Systems

Spring 2025

Yue Cheng



Midterm exam

- Thursday, February 27, 2:00 pm – 3:15 pm
 - Open book, open notes
- Covering four topics from Lec 2 to Lec 8
 - CPU scheduling policies
 - Caching policies
 - Hadoop MapReduce + HDFS
 - Spark RDD

Logistics

- The exam will be remote + synchronous over **gradescope**
- The exam sheet will be available on **gradescope** at 2 pm
- You should work directly on **gradescope**
- Submission closes at 3:25 pm (a grace period of 10 minutes for submission)

CPU job scheduling

- FIFO
 - How it works?
 - FIFO's problems (why we need SJF)?
- SJF
 - How it works?
 - Any limitations (why we need STCF)?
- STCF (preemptive SJF)
 - How it works? How it solves SJF's limitations?
- RR (Round Robin)
 - How it works?

Caching policy

- LRU (least recently used)
- FIFO (first-in, first-out)

MapReduce + HDFS

- How MapReduce works
- The performance characteristics of different phases of a MapReduce job (TeraSort)
- Fault tolerance
 - Storage level: Replication for HDFS
 - Compute level: Backup tasks for MapReduce

Spark

- Motivation
- Transformations and actions
 - Narrow vs. wide transformation
- `.cache()` to pin a computed RDD into memory to avoid recomputation

Question types

- Multi-choice questions (~40%)
- True or false questions (~30%)
- Problem solving (~30%)